

Something from nothing

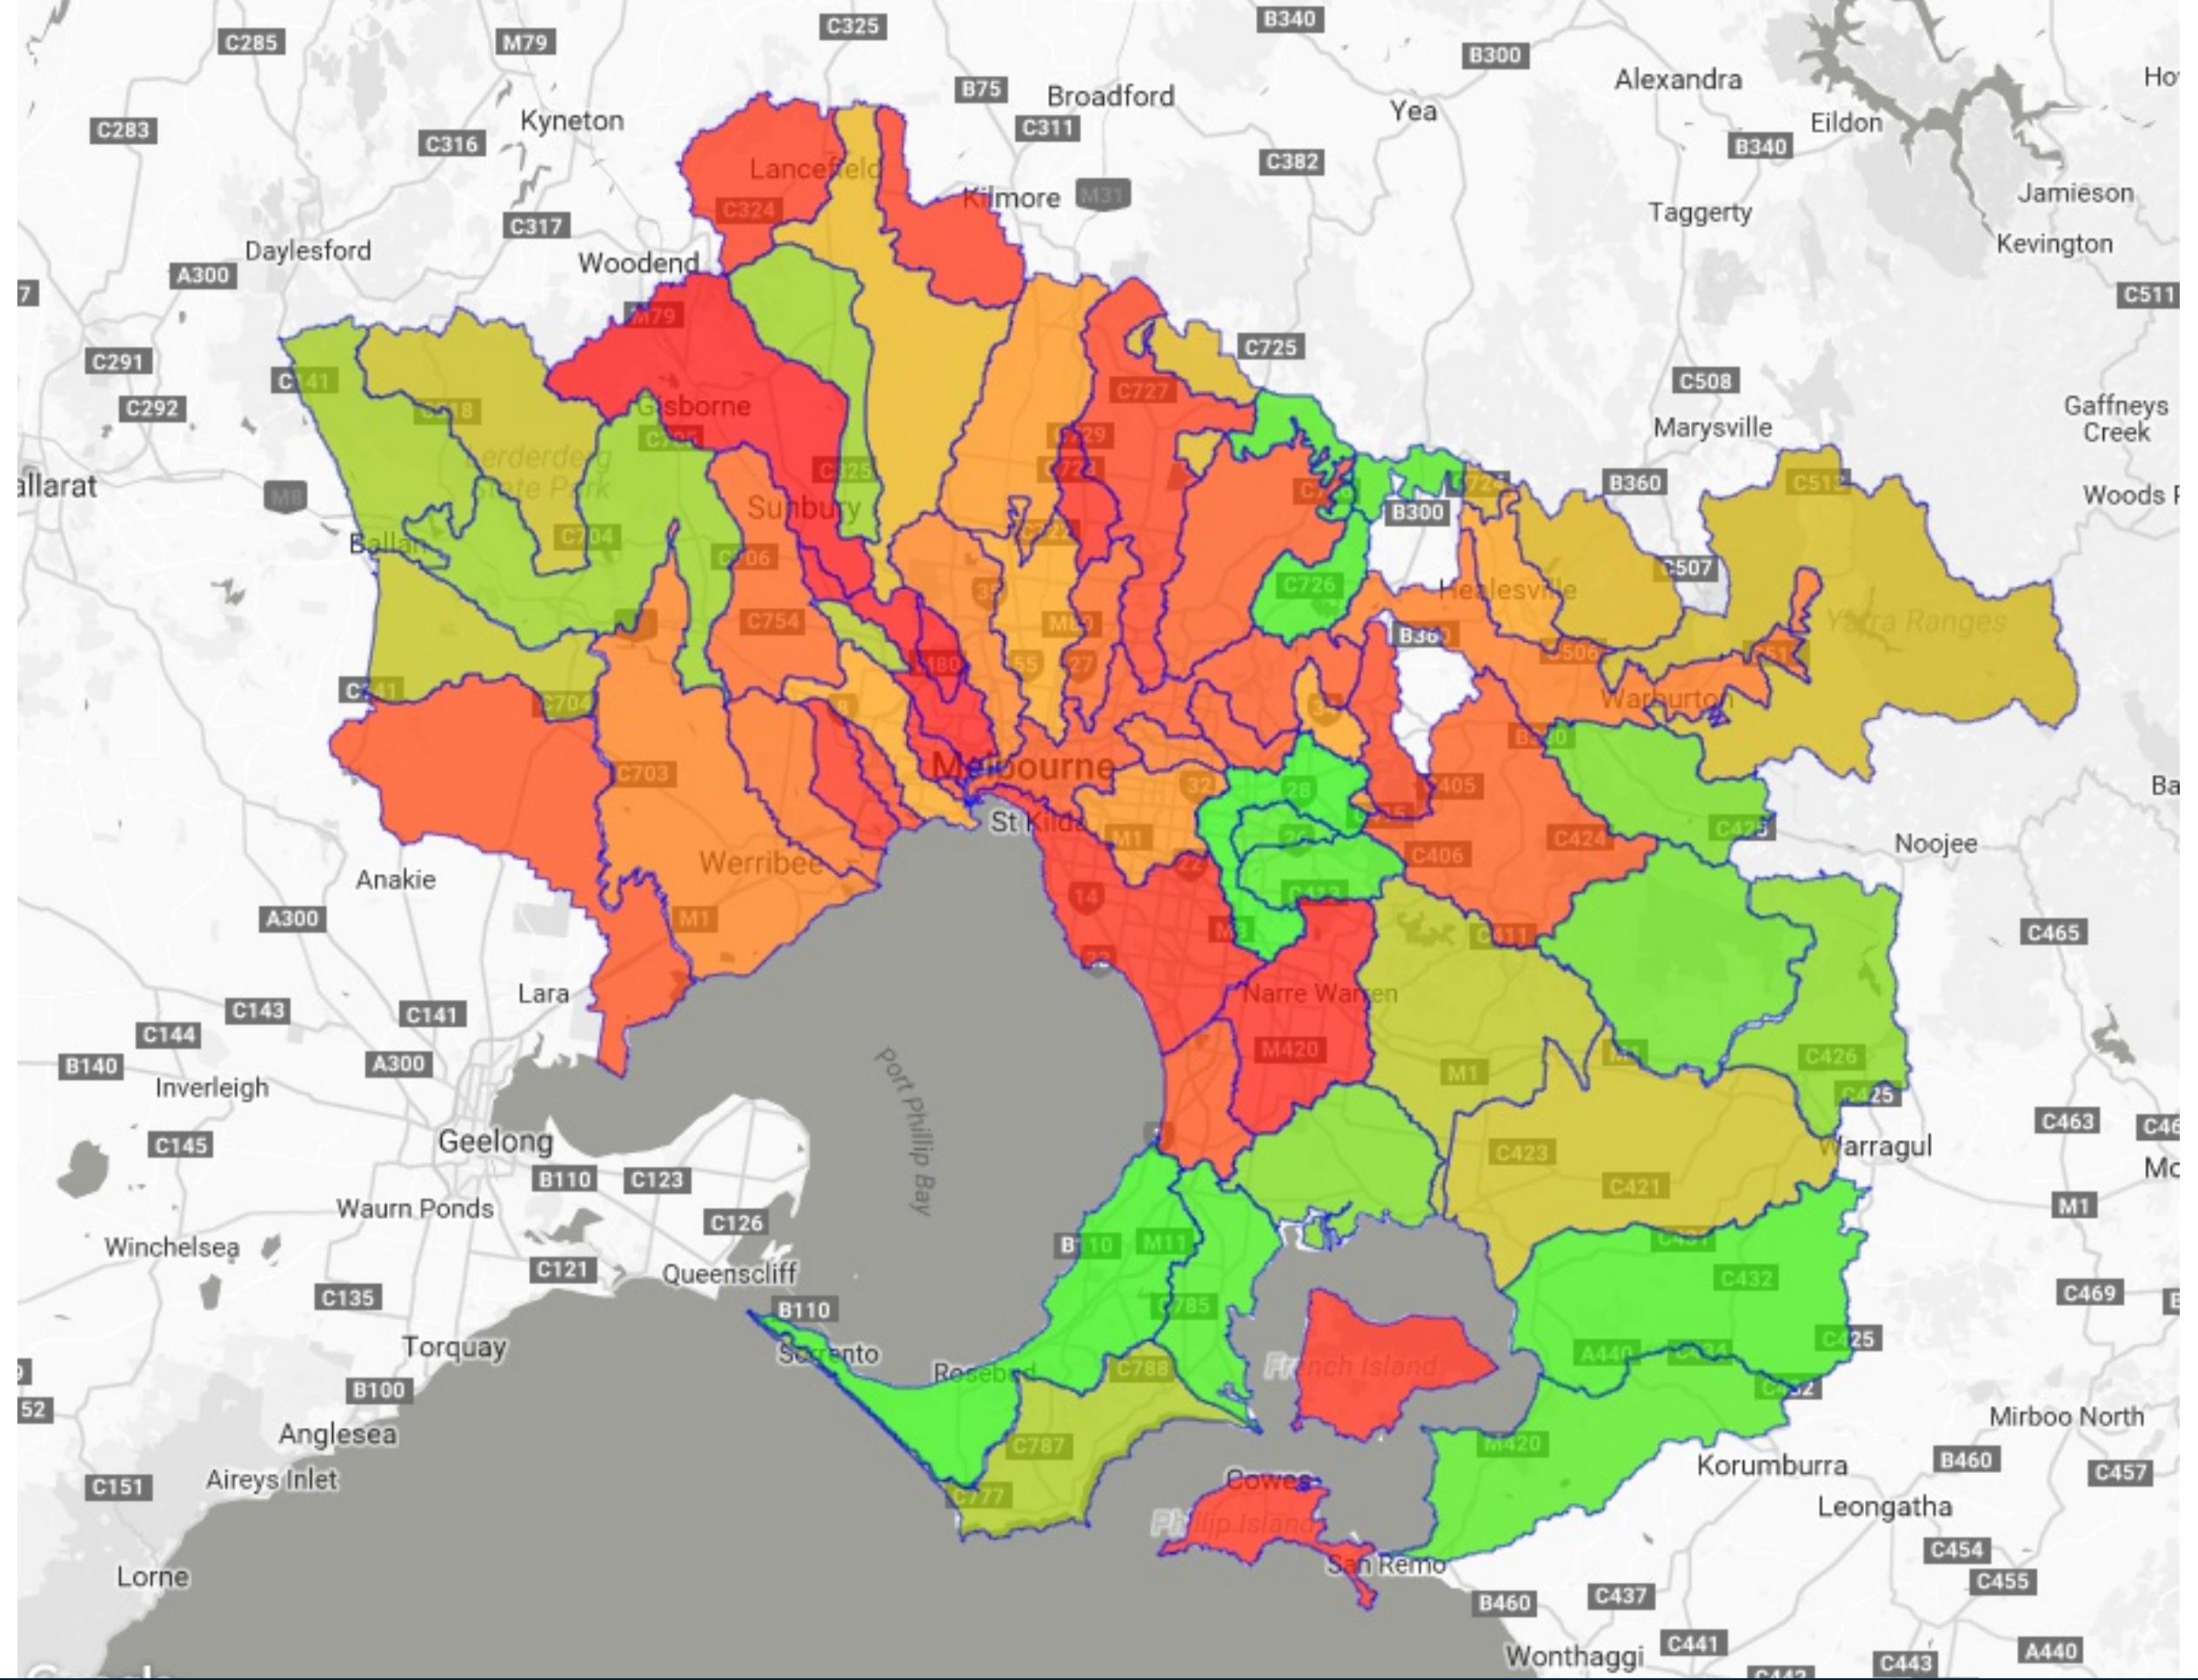
TensorFlow models from scratch

Elizabeth Stark

www.symbolix.com.au

It started with a frog ...





Frog Census



**Melbourne
Water**

The first challenge



Solution:



The second challenge



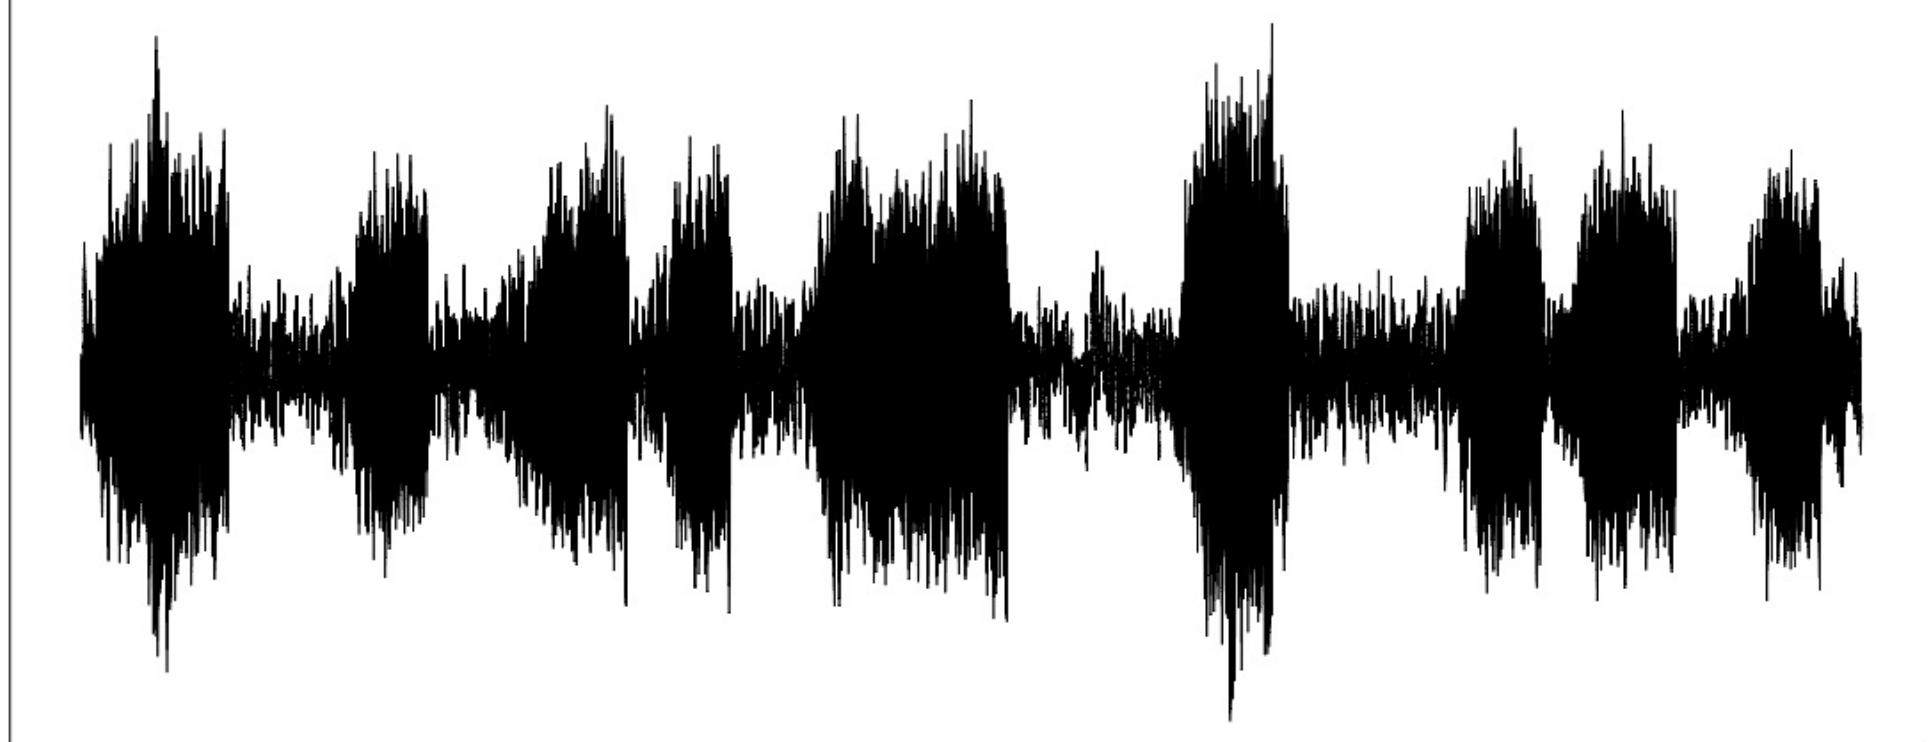
Southern Toadlet



Peron's Tree Frog



GGF



Computer listening vs computer vision



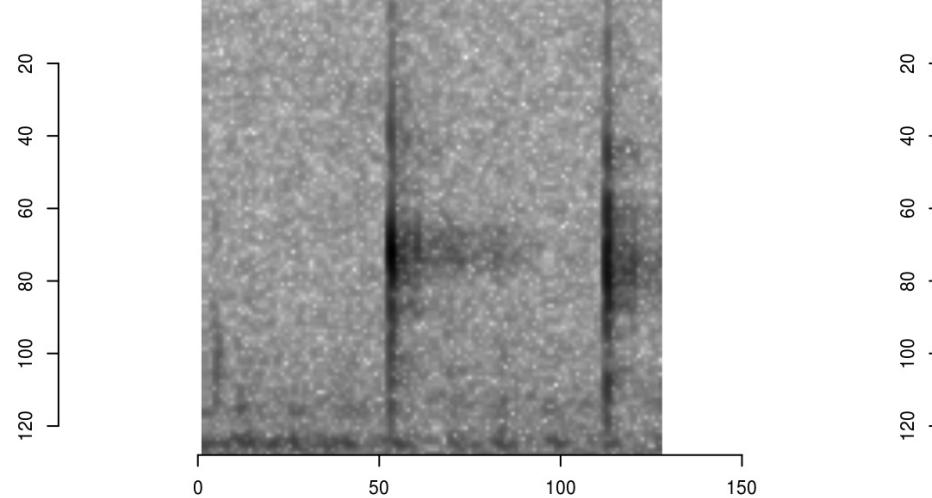
Keyword="sound" : ~100K papers

Keyword="image" : ~3.5Mil papers

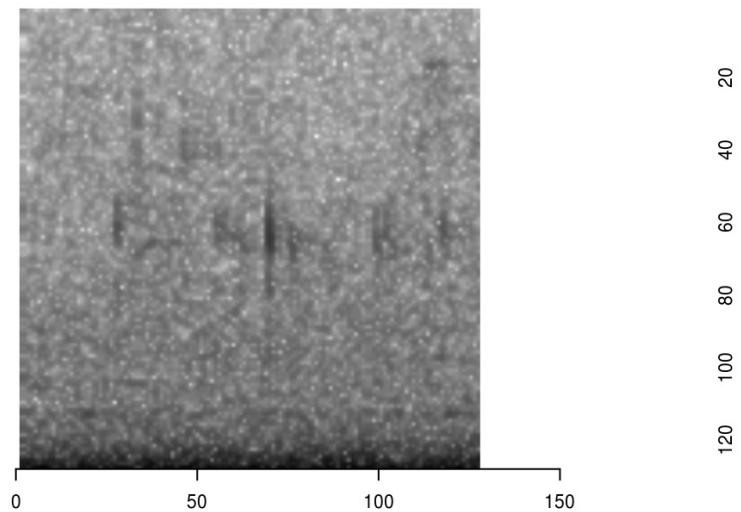
Solution:

Play with the cool kids toys computer
vision tools

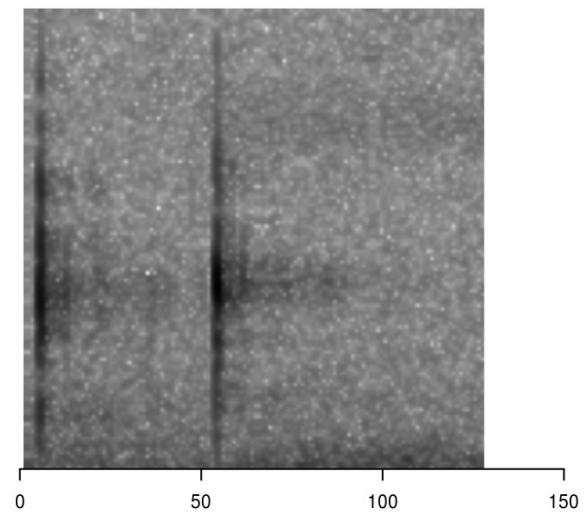
1



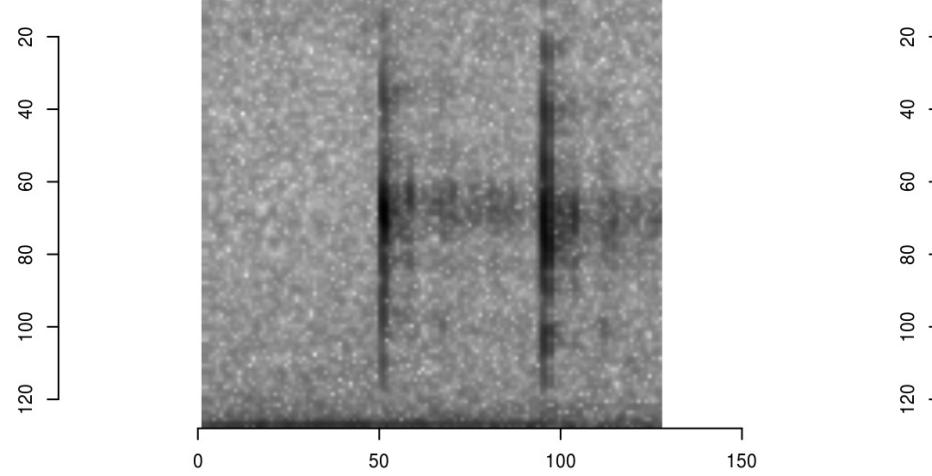
1



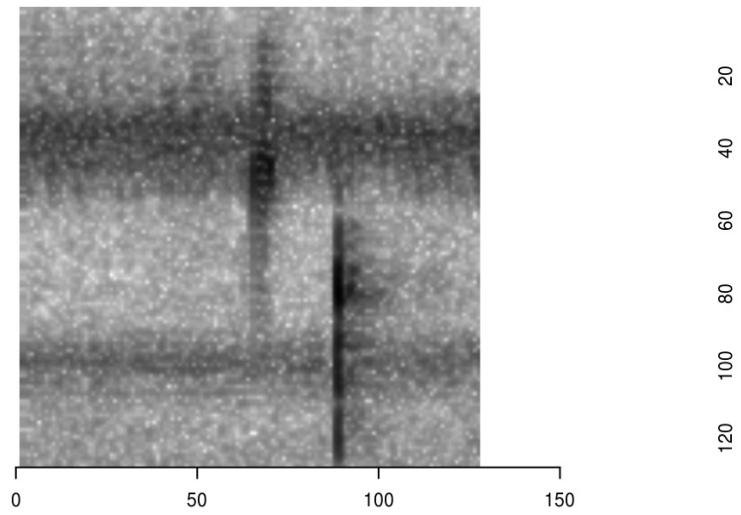
1



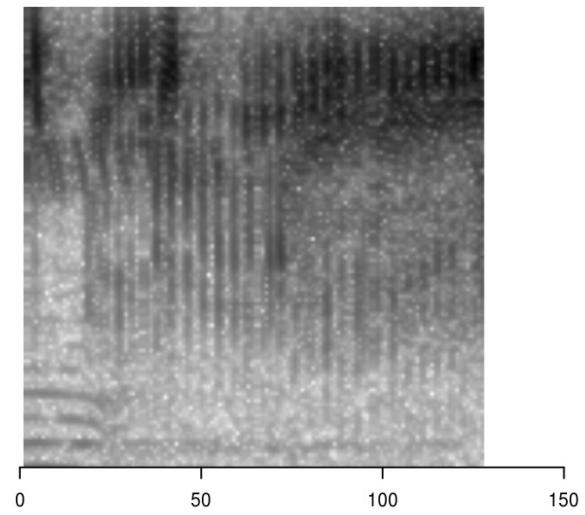
1



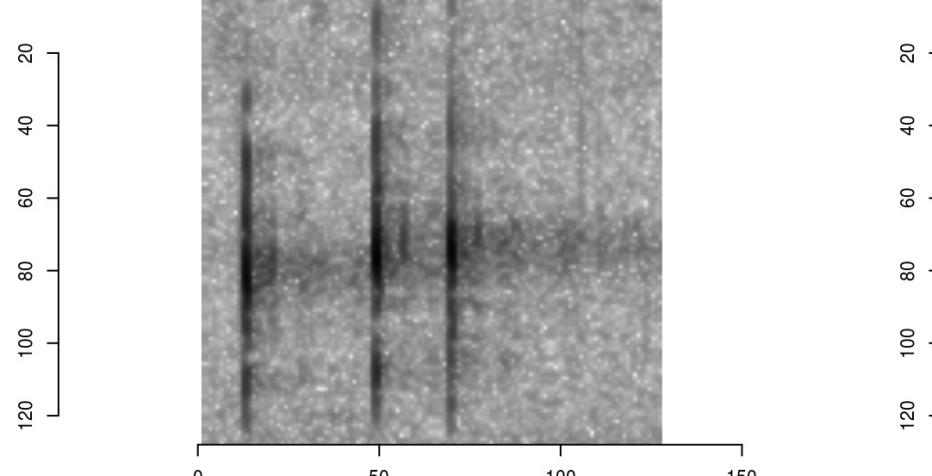
2



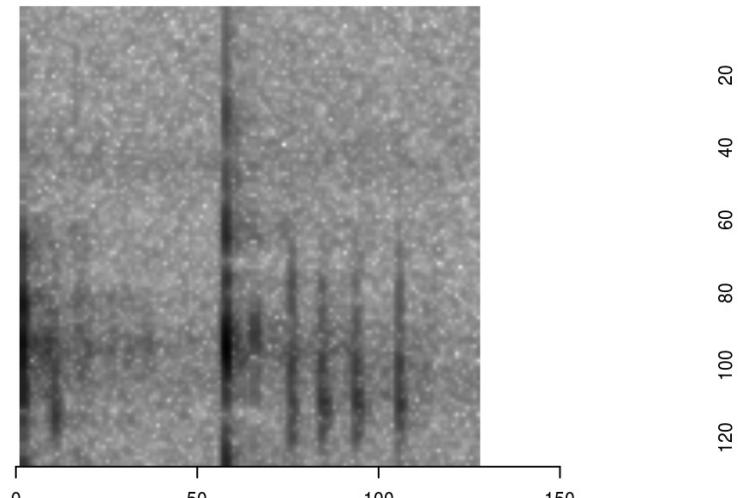
0



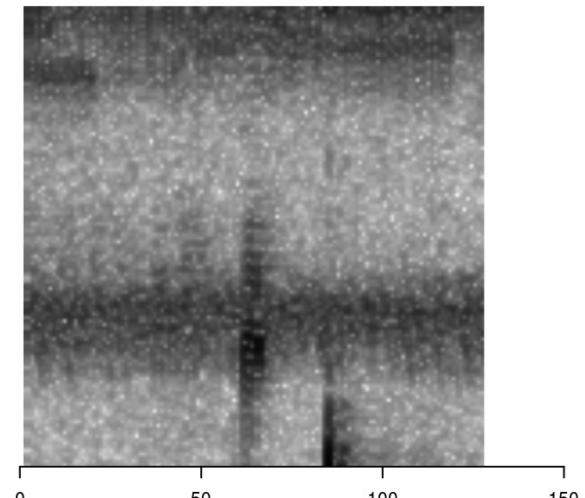
1



1



2



The third challenge



How few data?

- ~1,400 individual tagged calls
- 5 classes
 - *Litoria fallax* (Eastern Dwarf Tree frog)
 - *Limnodynastes tasmaniensis* (Spotted grass frog)
 - *Crinia signifera* (Common Eastern froglet)
 - *Limnodynastes peronii* (Striped Marsh frog)
 - Other frog

..... And about 100 million billion crickets calling over the top

Solution



Our toolbox

```

rs-pooling.R x layers-normalization.R x models.py x mnist_antirectifier.R x mnist_mlp.R x test-cus x
Source on Save Run Source Environment History Connections Build Git
24 x_train <- x_train / 255
25 x_test <- x_test / 255
26
27 cat(nrow(x_train), 'train samples\n')
28 cat(nrow(x_test), 'test samples\n')
29
30 # convert class vectors to binary class matrices
31 y_train <- to_categorical(y_train, num_classes)
32 y_test <- to_categorical(y_test, num_classes)
33
34 model <- keras_model_sequential()
35 model %>%
36   layer_dense(units = 256, activation = 'relu', input_shape = c(784)) %>%
37   layer_dropout(rate = 0.4) %>%
38   layer_dense(units = 128, activation = 'relu') %>%
39   layer_dropout(rate = 0.3) %>%
40   layer_dense(units = 10, activation = 'softmax')
41
42 summary(model)
43
44 model %>% compile(
45   loss = 'categorical_crossentropy',
46   optimizer = optimizer_rmsprop(),
47   metrics = c('accuracy')
48 )
49
50 history <- model %>% fit(
51   x_train, y_train,
52   batch_size = batch_size,
53   epochs = epochs,
54   verbose = 1,
55   validation_split = 0.2
56 )
57
45:37 (Top Level) ▾

```

Console Terminal Find in Files R Markdown

```

~/packages/keras/
+ )

> history <- model %>% fit(
+   x_train, y_train,
+   batch_size = batch_size,
+   epochs = epochs,
+   verbose = 1,
+   validation_split = 0.2
+ )
Train on 48000 samples, validate on 12000 samples
Epoch 1/30
48000/48000 [=====] - 2s - loss: 0.4333 - acc: 0.8675 - val_loss: 0.1725 - val_acc:
0.9495
Epoch 2/30
48000/48000 [=====] - 2s - loss: 0.2037 - acc: 0.9401 - val_loss: 0.1240 - val_acc:
0.9622
Epoch 3/30
48000/48000 [=====] - 2s - loss: 0.1602 - acc: 0.9531 - val_loss: 0.1065 - val_acc:
0.9690
Epoch 4/30
19072/48000 [=====>.....] - ETA: 1s - loss: 0.1324 - acc: 0.9610

```

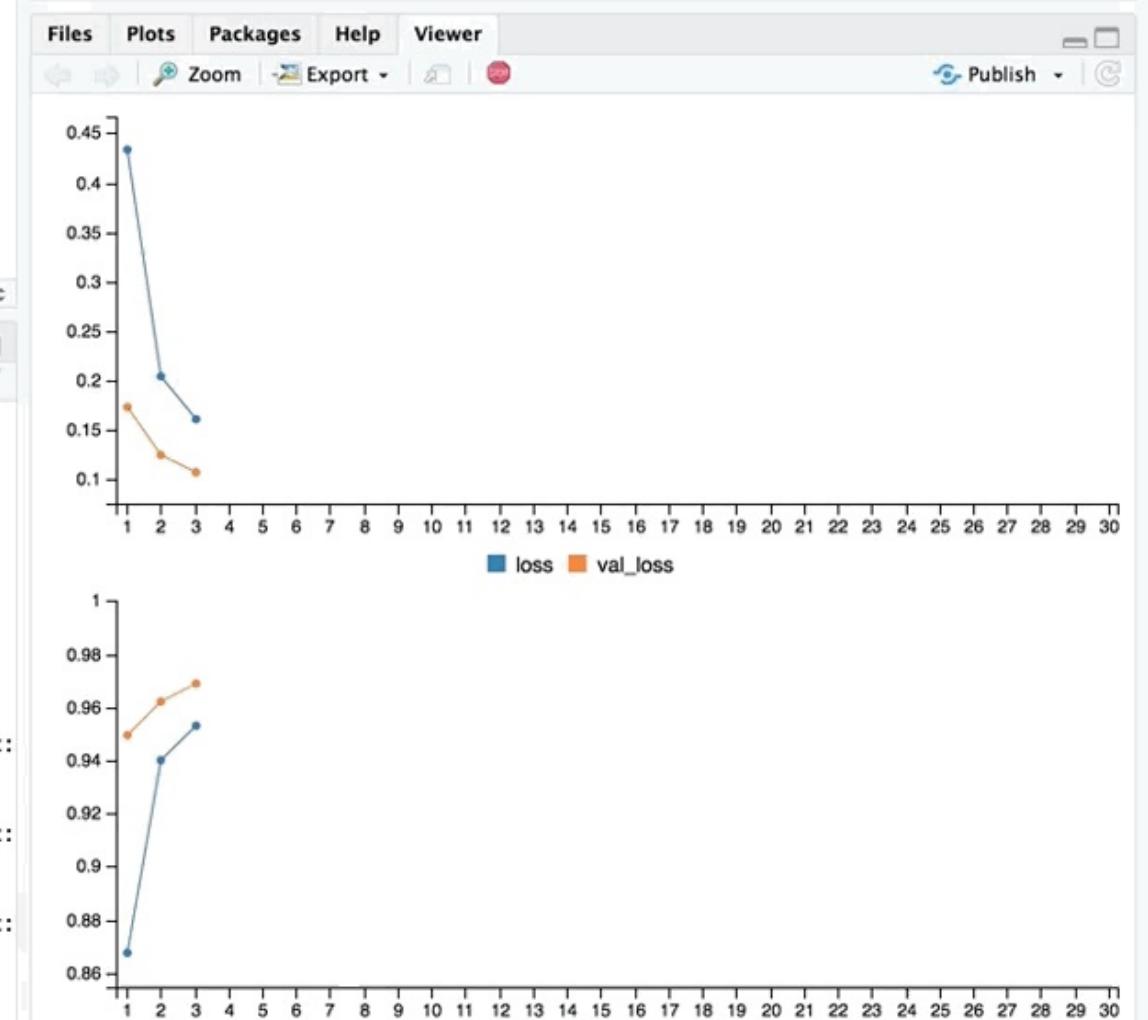
Import Dataset

Global Environment

history	List of 2
mnist	Large list (2 elements, 209.6 Mb)
score	List of 2
x_test	Large matrix (7840000 elements, 59.8 Mb)
x_train	Large matrix (47040000 elements, 358.9 Mb)
y_test	Large matrix (100000 elements, 781.4 Kb)
y_train	Large matrix (600000 elements, 4.6 Mb)

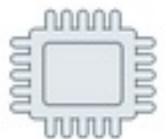
Values

batch_size	128
epochs	30
model	Model
num_classes	10



With all the usual tools

[GPUs](#)



[CloudML](#)



[Training Flags](#)



[Training Runs](#)



[TensorBoard](#)



[Datasets API](#)

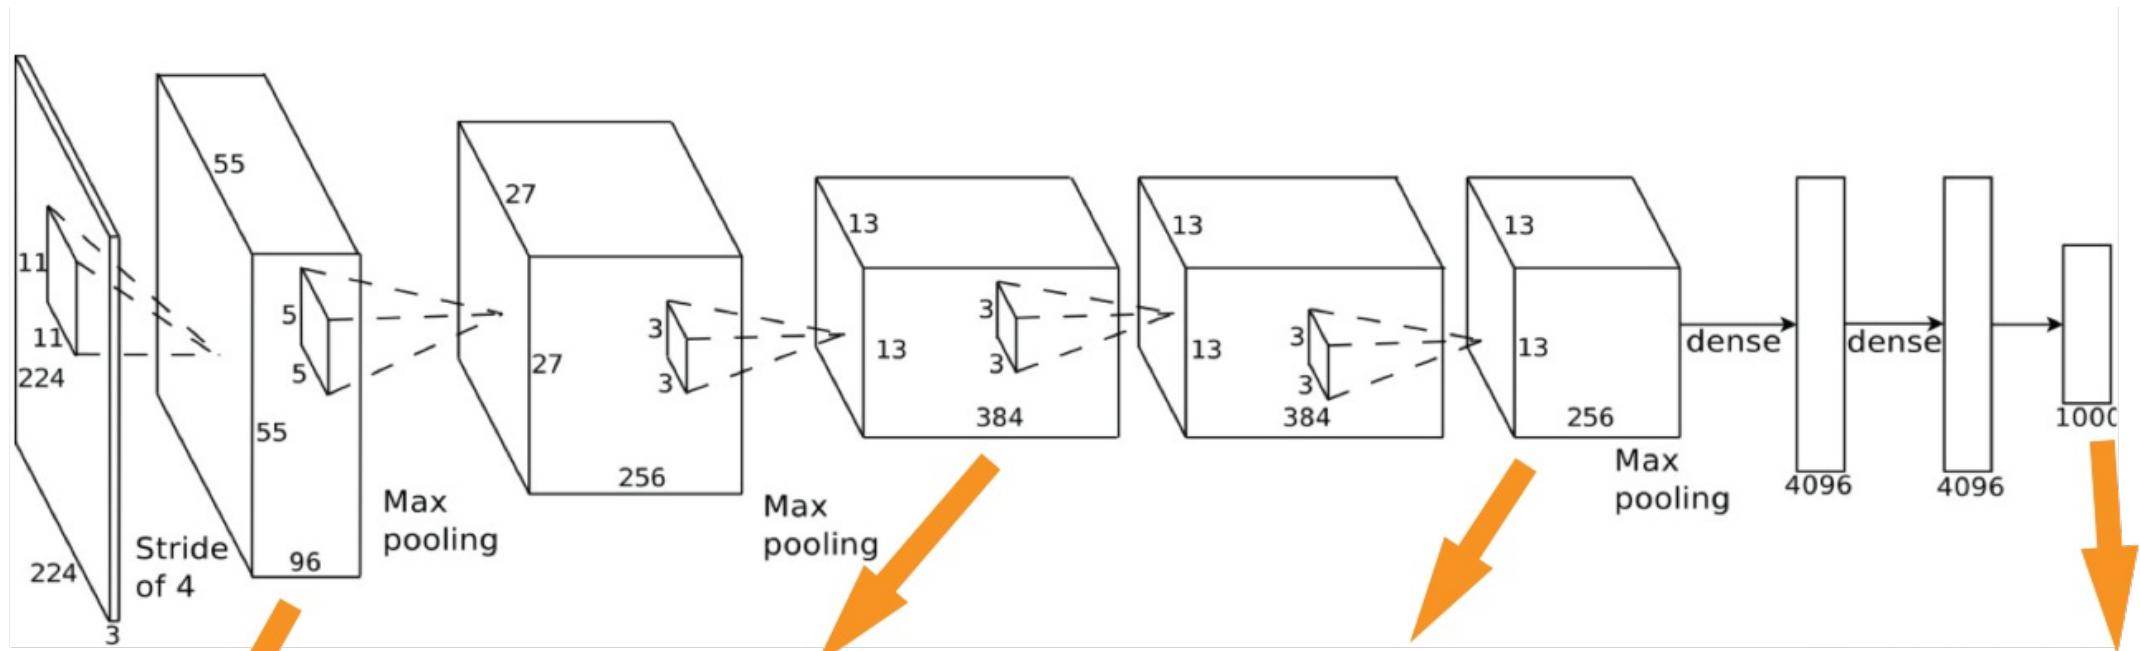


[Deployment](#)

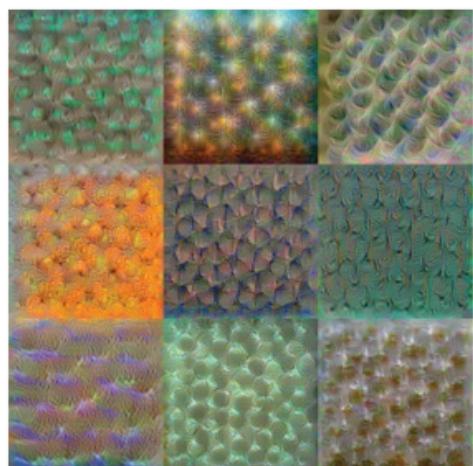


Tricks and tips

Leave the deepest networks for the big datasets



Conv 1: Edge+Blob



Conv 3: Texture

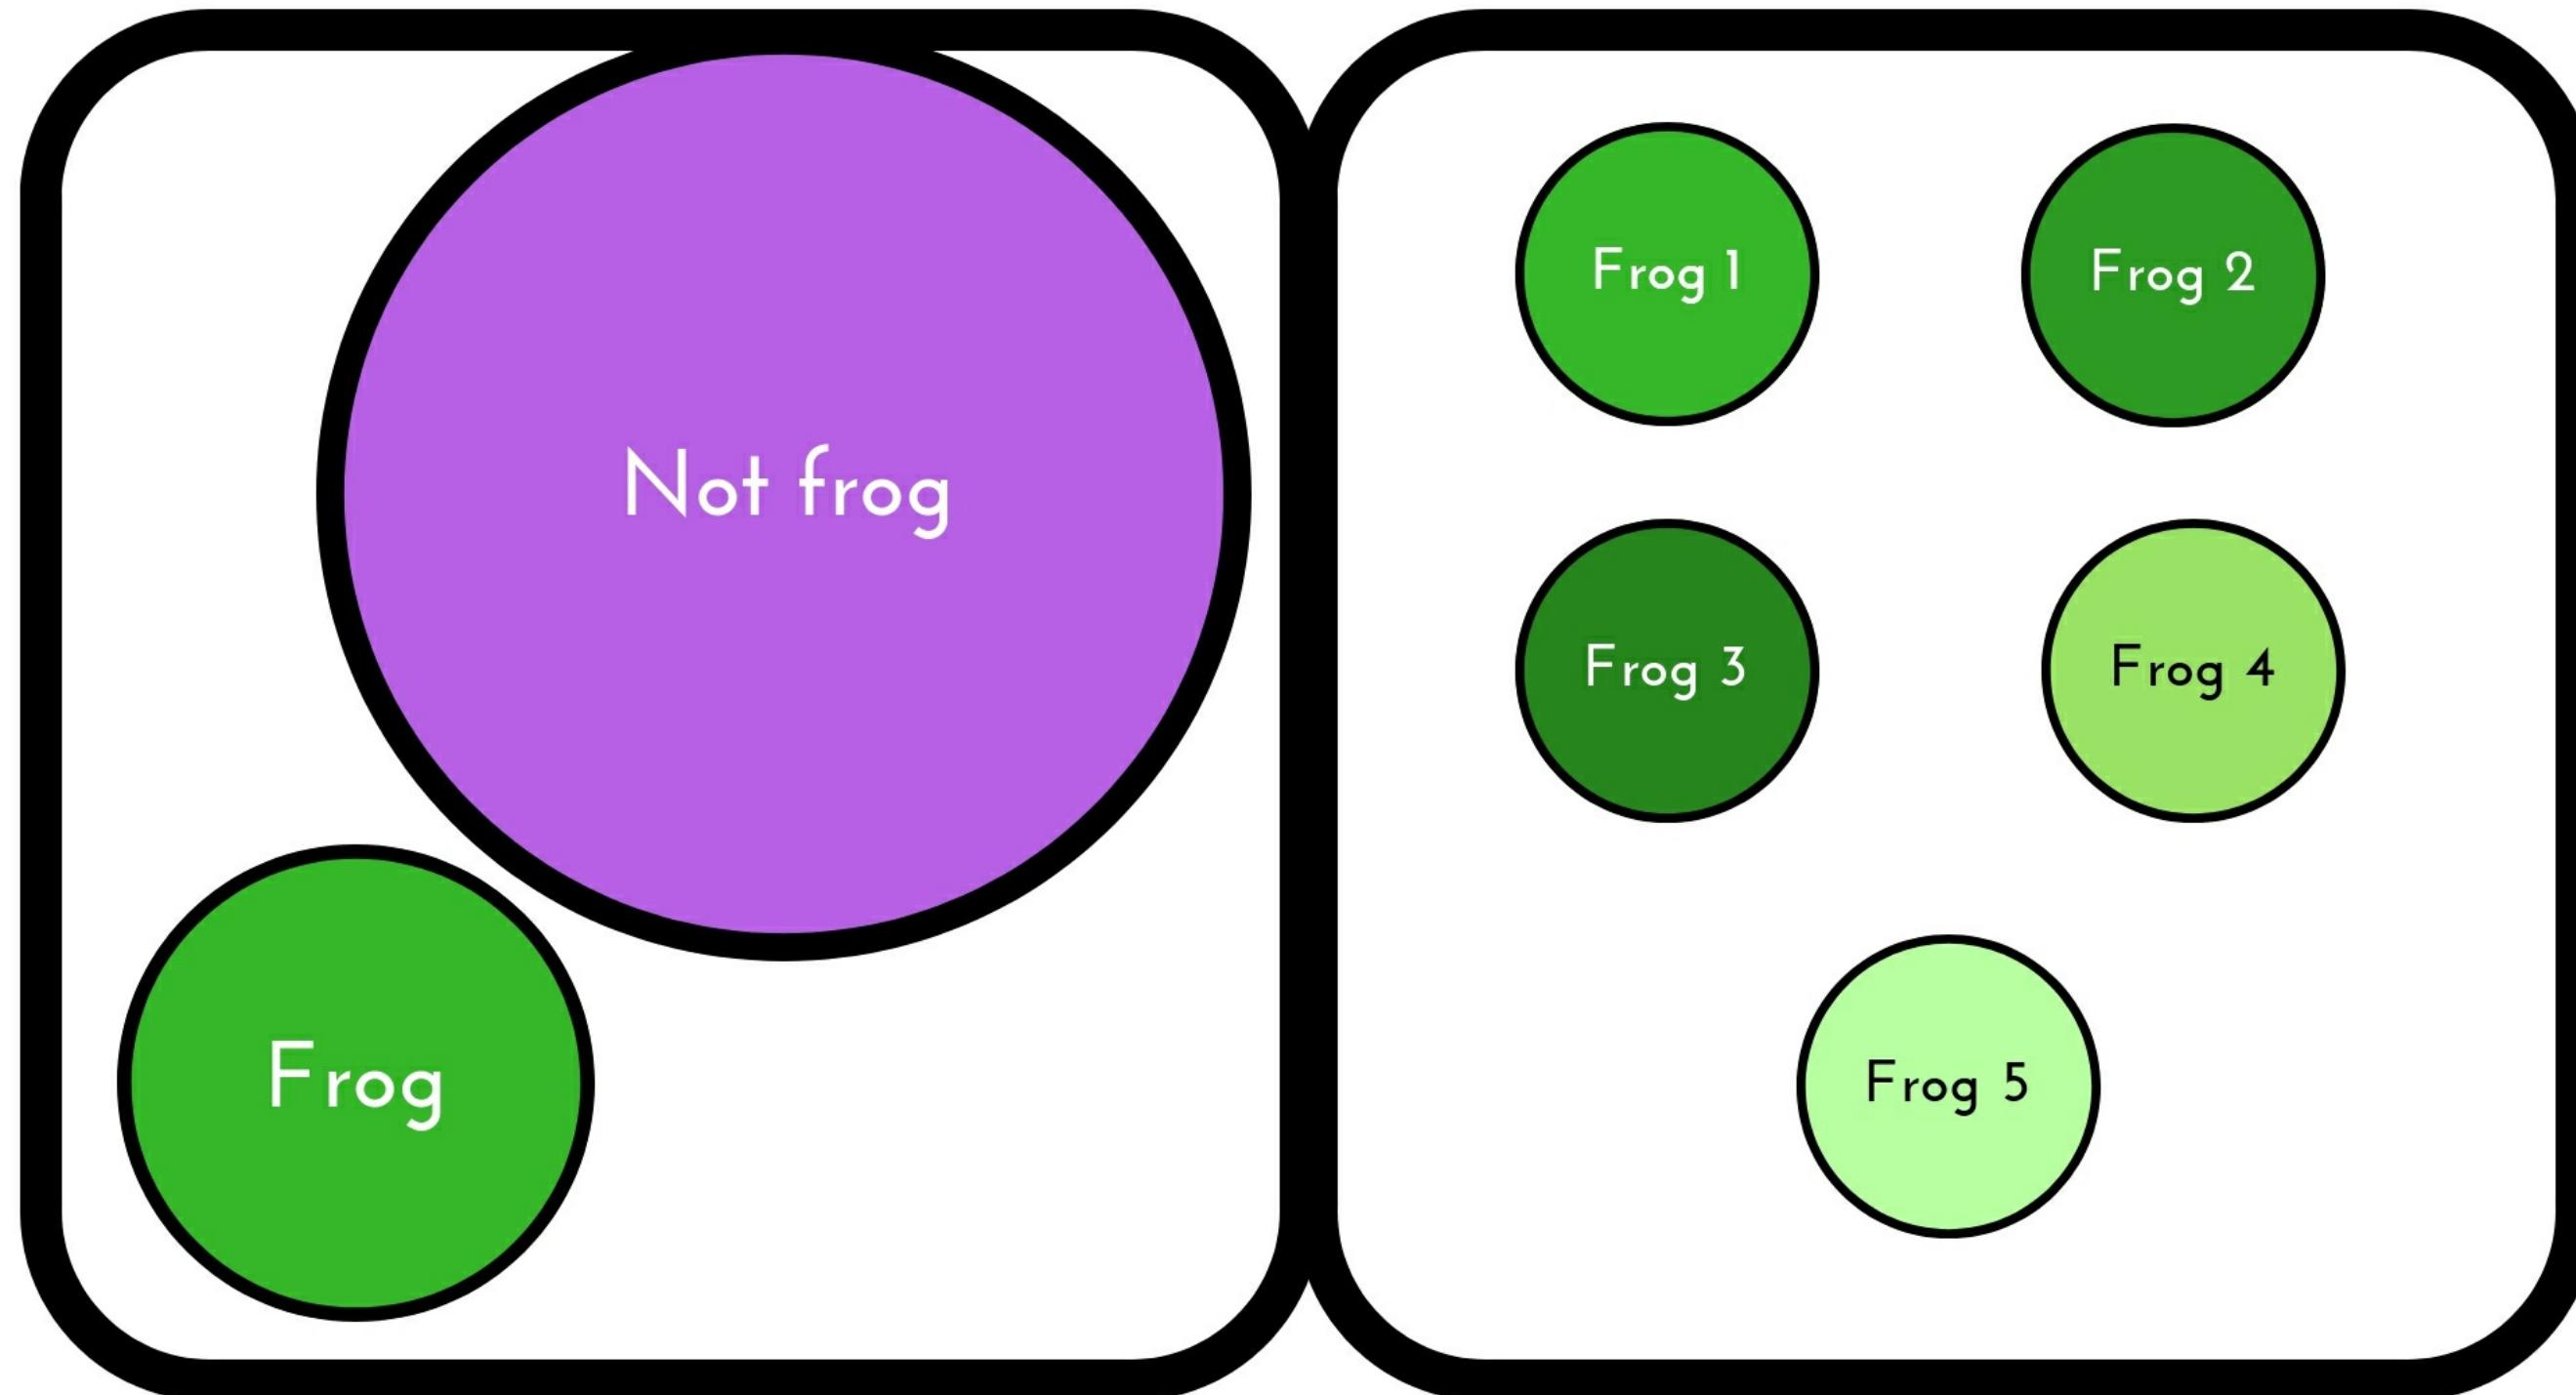


Conv 5: Object Parts



Fc8: Object Classes

Even it out

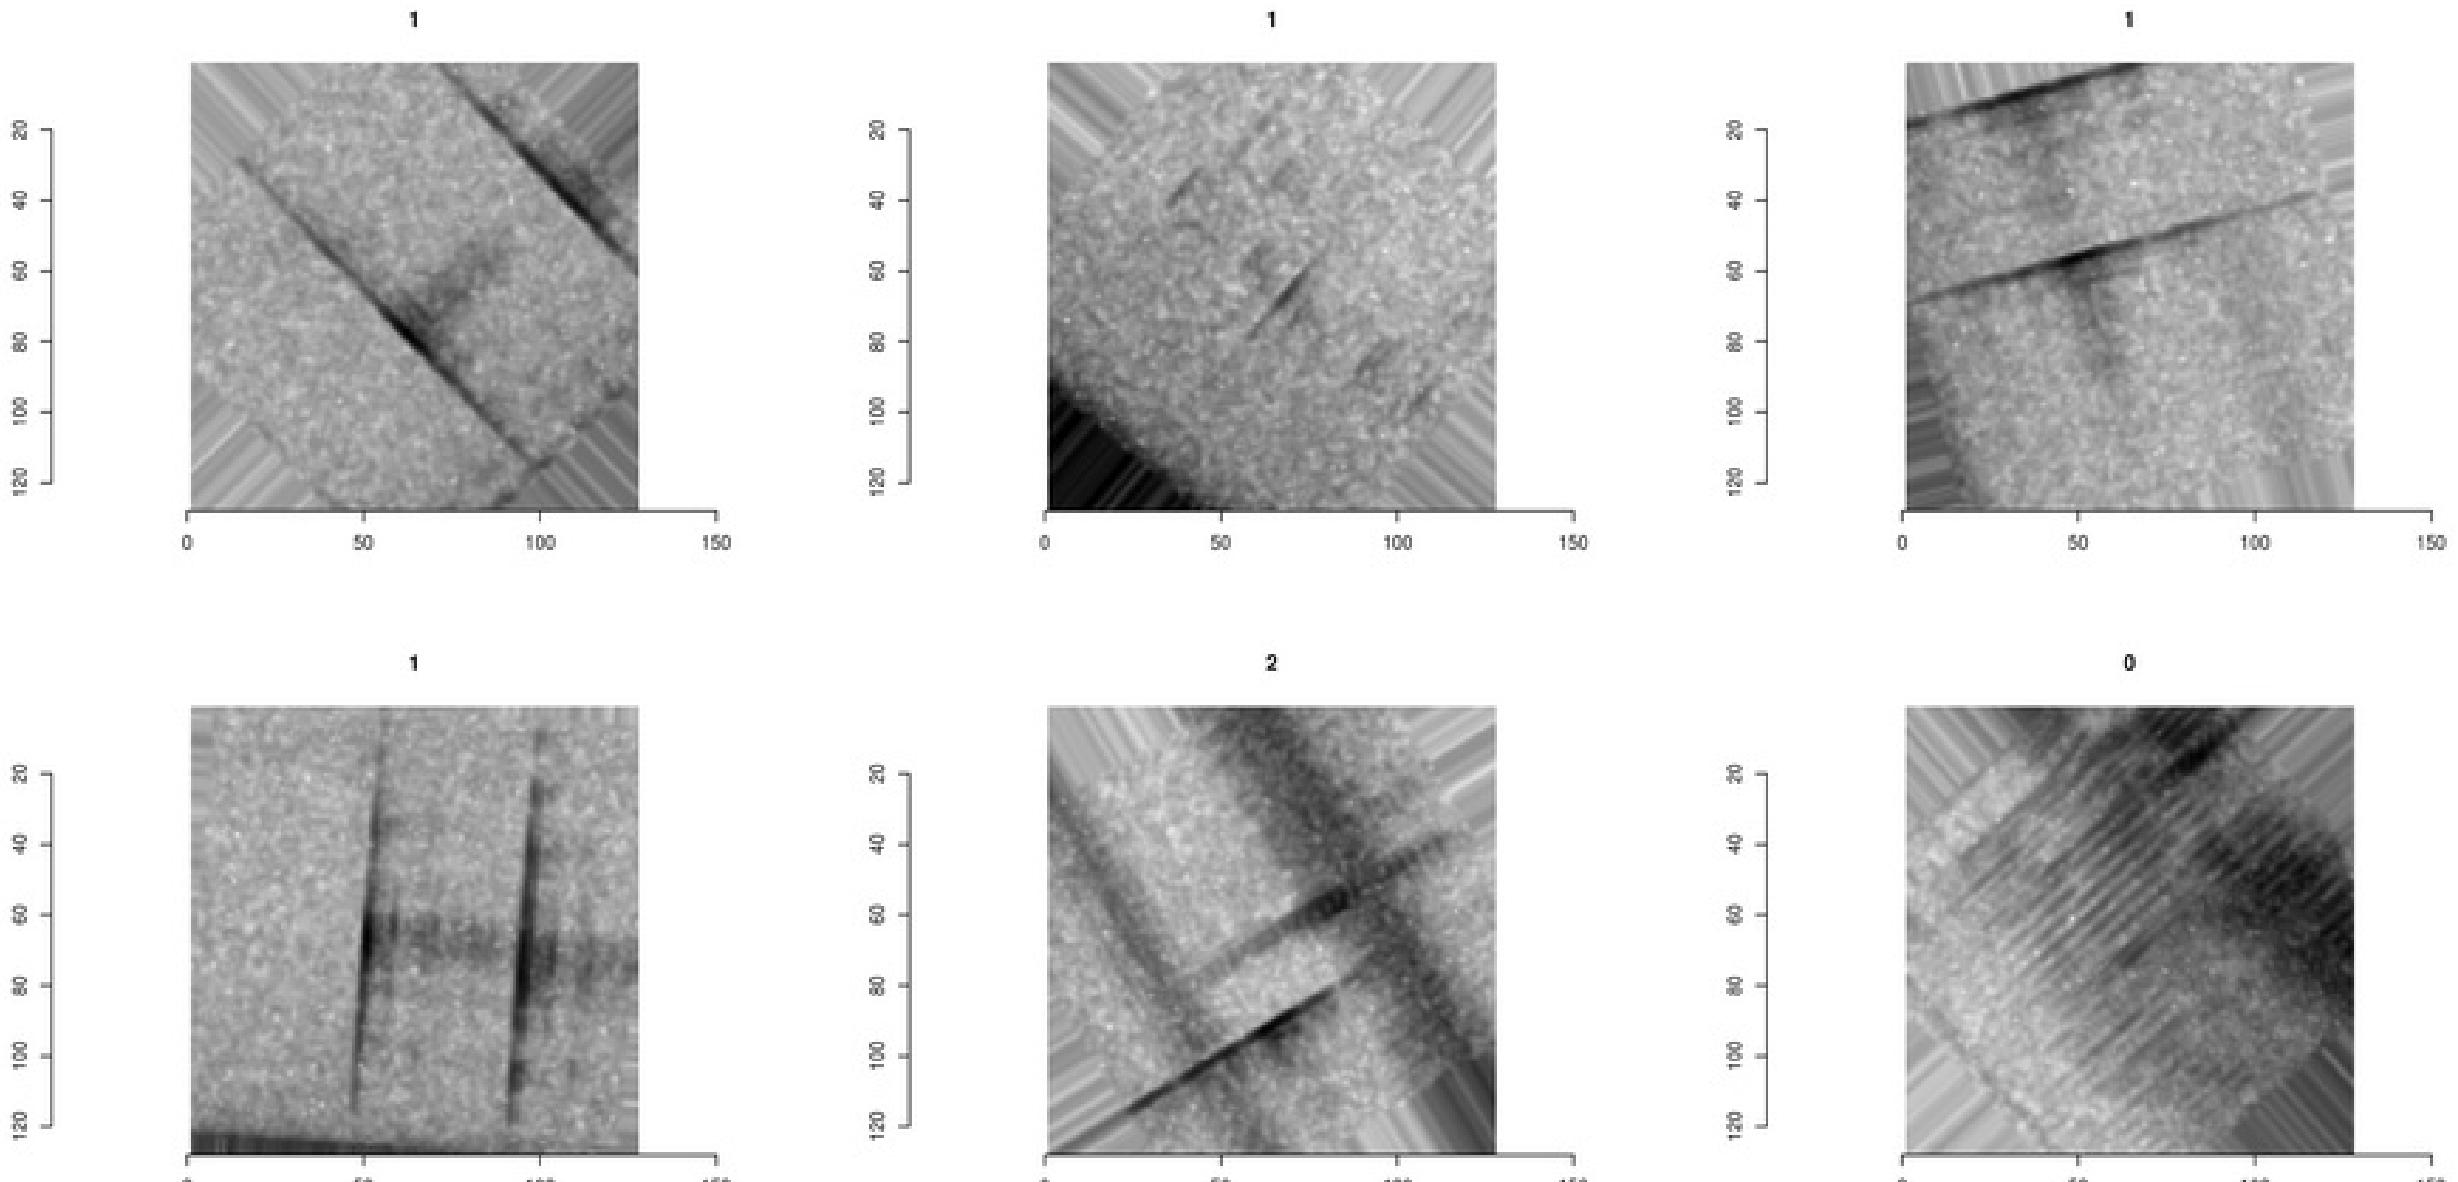


Data augmentation

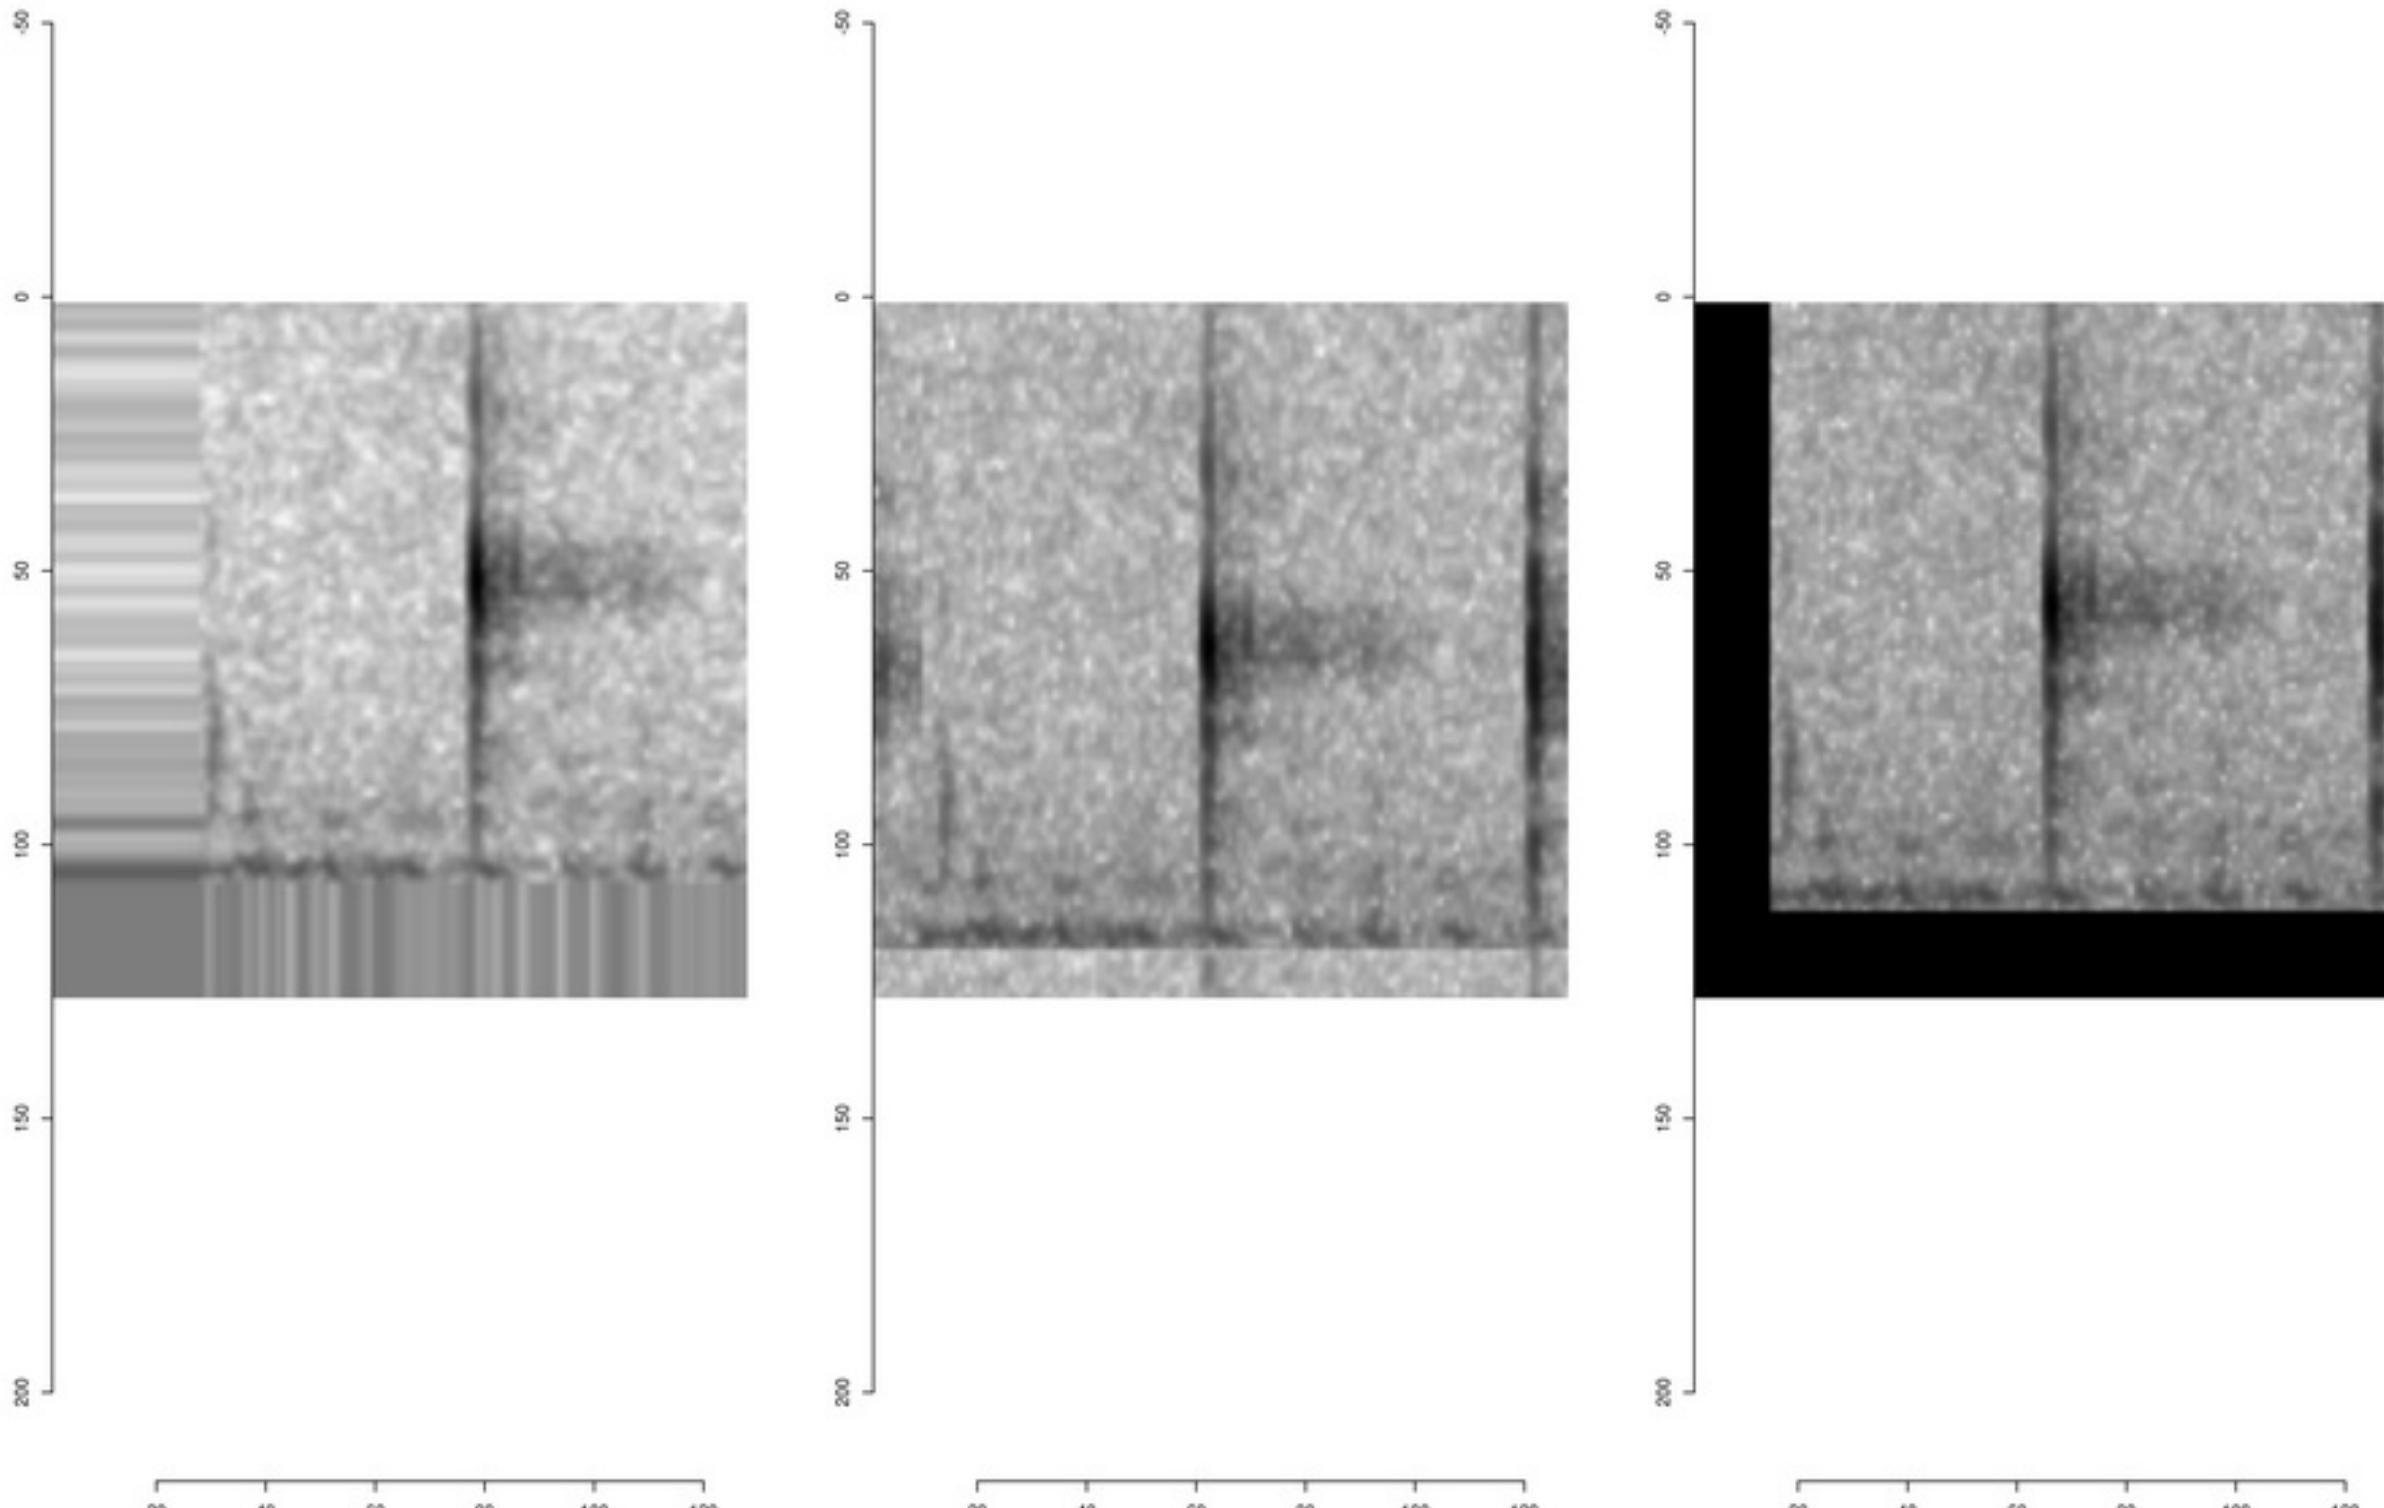
- Normalise/standardise
- Rotate
- Shift/Jitter
- Shear
- Zoom
- Flip

Rotation

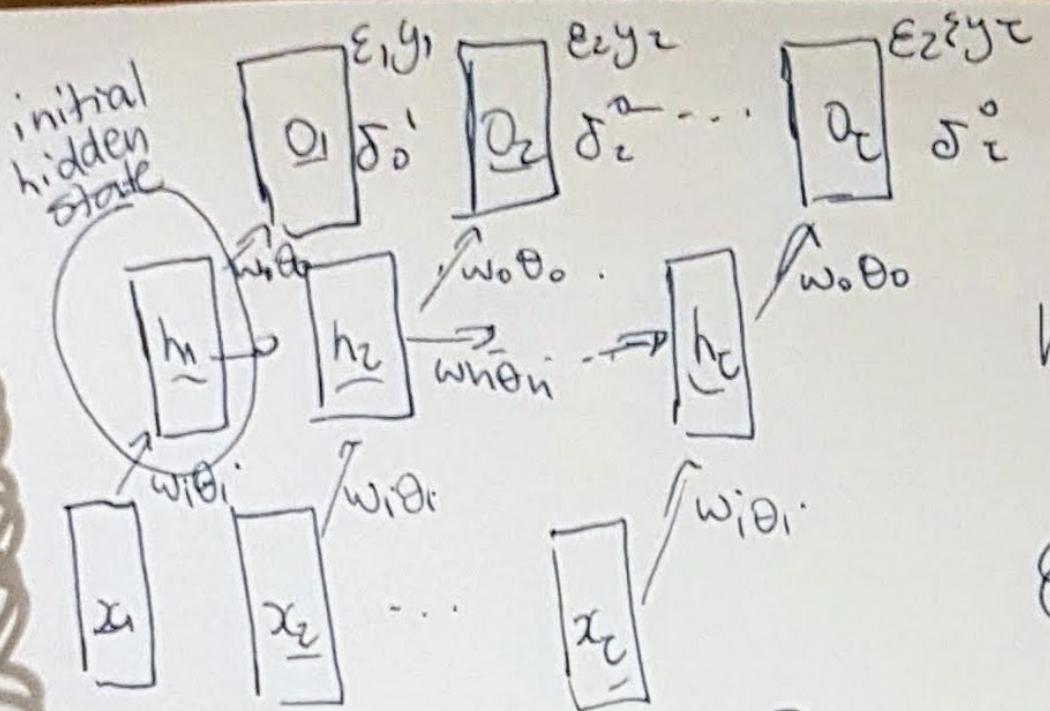
```
data_gen_aug <- image_data_generator(  
    rescale = 1,  
    featurewise_center = FALSE,  
    featurewise_std_normalization = FALSE,  
    samplewise_center = TRUE,  
    samplewise_std_normalization = TRUE,  
    rotation_range = 90    # in degrees  
)  
  
batch_aug_rotate <- augBatch(object = data_gen_aug, img = img, lbl = lbl)
```



Shift



**Sometimes you need the big
guns**



BACKPROP

: UNROLL THE NETWORK
FOR τ LENGTH

$$x = \langle \underline{x_1} \underline{x_2} \dots \underline{x_\tau} \rangle$$

$$h_0 = \underbrace{(0, 0, 0, \dots)}_{\dots}$$

$$\epsilon = \sum_{i=1}^{\tau} \epsilon_i \rightarrow \text{just add errors.}$$

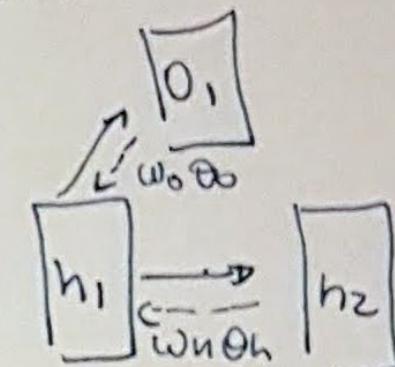
$$\text{eg } \epsilon_i = \|y_i - o_i\|^2$$

compute the δ 's

$$\delta_t^o = \partial f^o / \partial \epsilon_t \quad \forall t \in [1 \dots \tau]$$

$$\delta_t^h = \partial f^h / \partial \left(\underbrace{w_0 \delta_t^o + w_n \delta_{t+1}^h}_{\text{backprop from output}} \right)$$

$$\delta_t^h = \partial f^h / \partial \left(\underbrace{w_0 \delta_t^o}_{\text{from next hidden state } t+1} + \underbrace{w_n \delta_{t+1}^h}_{\text{from next hidden state } t+1} \right)$$



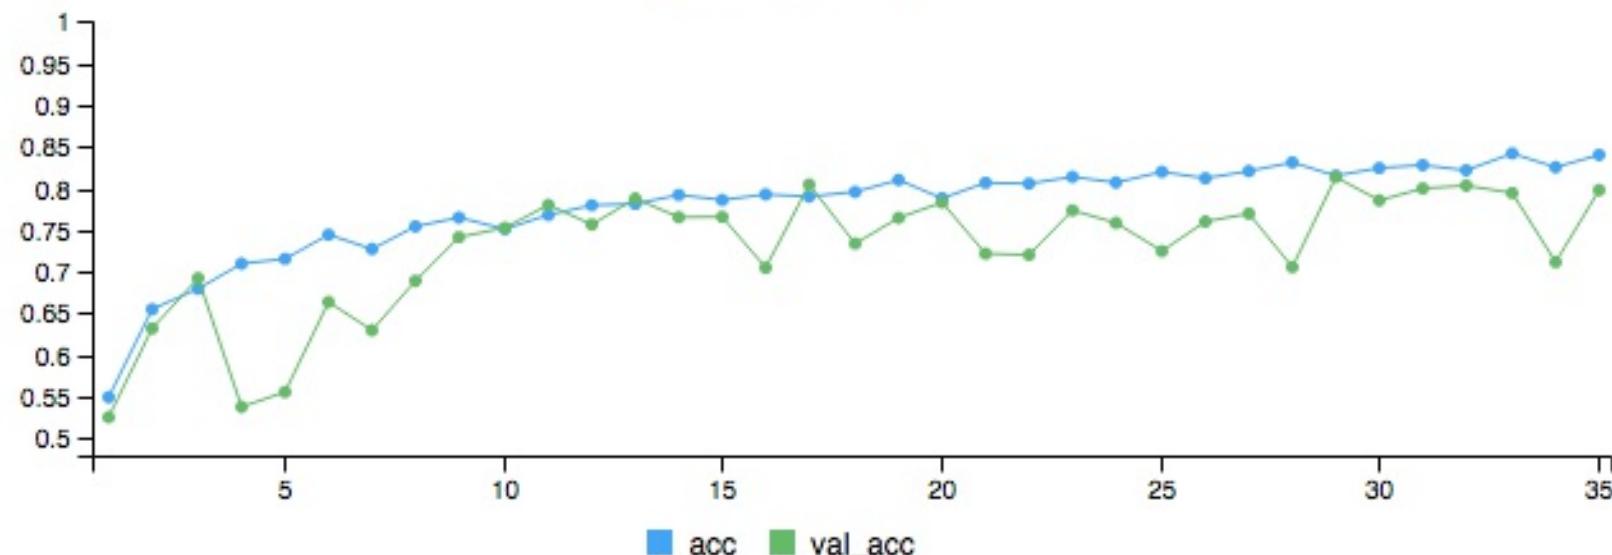
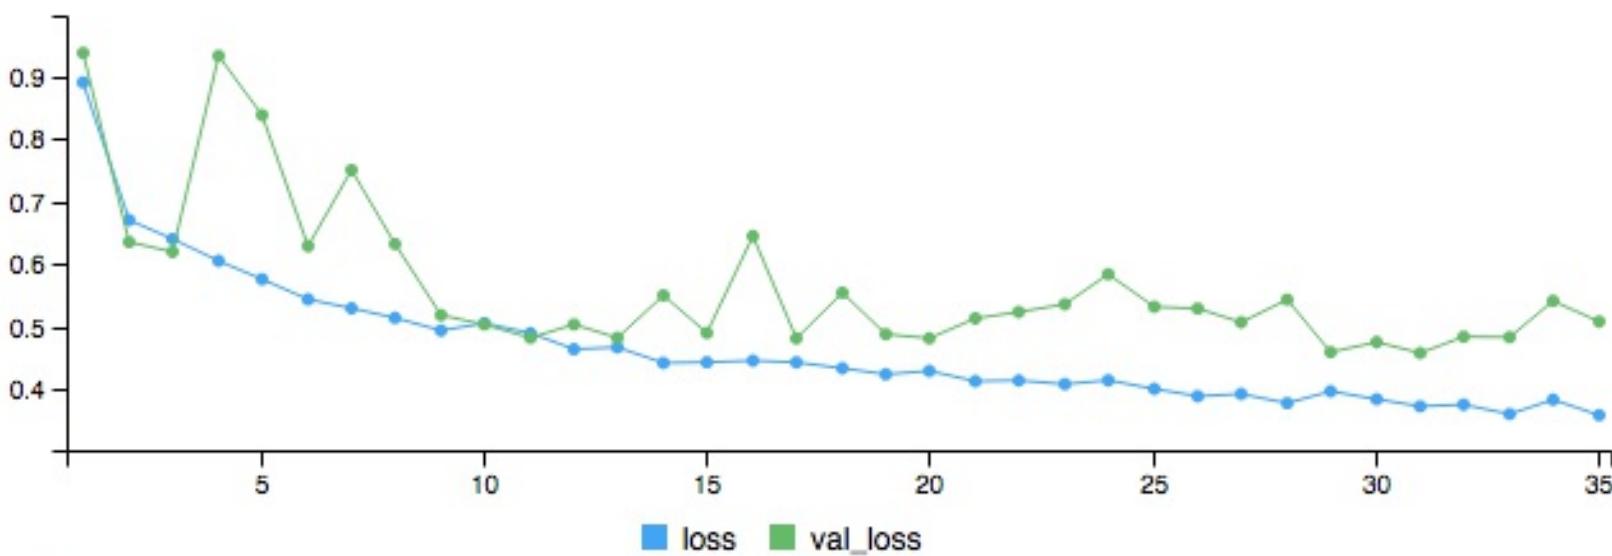
then

$$\nabla w_0 = \sum_{t=1}^{\tau} h_t \cdot (\delta_t^o)^2$$

Bidirectional ~~LSTMs~~. RNNs

The result?

History



Run

context	local
script	Conv4Dense2_asymconv.R
started	2018-05-08 02:57:56 GMT
time	00:11:52

Metrics

loss	0.3590
acc	0.8403
val_loss	0.5083
val_acc	0.7981

Flags

model_clayers	3
conv_activation	relu
dense_activation	relu
ksz1	5
ksz2	3
mp1	2

Next steps



VectorStock®

VectorStock.com/988530



